

---

# Deep Security Analysis of NVIDIA NeMo Ecosystem: From Probabilistic Generation to Deterministic Control

---

Anonymous Authors<sup>1</sup>

## Abstract

In the era of large-scale enterprise deployment of generative AI systems, the fundamental challenge for security architects is no longer merely the inference capability of foundation models, but rather how to construct a deterministic control layer atop an inherently probabilistic system. Large Language Models (LLMs), as stochastic parrots, exhibit unpredictable outputs that fundamentally conflict with enterprise requirements for compliance, security, and rigorous business logic. This paper presents a comprehensive security analysis of the NVIDIA NeMo ecosystem, with particular focus on NeMo Guardrails, the Colang modeling language, and NVIDIA Inference Microservices (NIM) integration architecture. Our analysis demonstrates that the NeMo ecosystem represents the most mature attempt in the industry for "programmable dialog management," introducing an event-driven state machine runtime that transforms unstructured natural language interactions into structured, executable flows through Colang. This architectural design enables security policies to be dynamic, context-aware interaction logic rather than static rule lists. We argue that NeMo Guardrails effectively establishes a deterministic bridge between infinite natural language expressions and finite security policies through canonical forms, enables asynchronous parallel defense through Colang 2.0's concurrent runtime, and provides zero-trust infrastructure through NIM's containerized security mechanisms. Our evaluation covers five categories of guardrails (input, dialog, retrieval, execution, and output), advanced RAG security mechanisms, data curation with NeMo Curator, and performance optimization strategies to minimize the "latency tax" of security enforcement.

## 1. Introduction

The deployment of Large Language Models (LLMs) in enterprise environments presents a unique security challenge: how to enforce deterministic security policies on systems that are fundamentally probabilistic. Traditional security mechanisms, designed for deterministic systems, are insufficient for LLM-based applications where outputs are generated stochastically and can vary significantly across identical inputs.

NVIDIA NeMo represents a comprehensive ecosystem addressing this challenge through a multi-layered security architecture. At its core, NeMo Guardrails provides a programmable framework for enforcing security policies through an event-driven runtime, while Colang serves as a domain-specific language for expressing security logic. The integration with NVIDIA Inference Microservices (NIM) ensures infrastructure-level security through containerized, signed model deployments.

This paper makes the following contributions:

- We provide the first comprehensive security analysis of the NVIDIA NeMo ecosystem, examining its architecture from both theoretical and practical perspectives.
- We analyze the evolution from Colang 1.0 to 2.0, demonstrating how asynchronous concurrency addresses performance bottlenecks in multi-modal security detection.
- We evaluate the effectiveness of canonical forms as a mechanism for bridging infinite natural language expressions to finite security policies.
- We present a detailed analysis of the five-category guardrail system and its application in enterprise scenarios.
- We examine RAG security mechanisms, data curation strategies, and performance optimization techniques specific to LLM guardrails.

---

<sup>1</sup>Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. **AUTHORERR: Missing \icmlcorrespondingauthor.**

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

## 055 2. Background and Related Work

### 056 2.1. LLM Security Challenges

057  
058 The security challenges of LLM deployment have been  
059 extensively documented. Prompt injection attacks (Goodside,  
060 2021) exploit the model’s inability to distinguish between  
061 user instructions and system instructions. Jailbreaking tech-  
062 niques (Wei et al., 2023) attempt to bypass safety filters  
063 through adversarial prompting. Hallucination (Ji et al.,  
064 2023) introduces factual errors that can lead to misinfor-  
065 mation or compliance violations.

066 Traditional approaches to LLM security have relied primarily  
067 on prompt engineering (White et al., 2023) or post-hoc  
068 filtering. However, these methods suffer from limited  
069 expressiveness and high false positive rates. More sophis-  
070 ticated approaches include fine-tuning models for safety  
071 (Ouyang et al., 2022) and reinforcement learning from hu-  
072 man feedback (RLHF) (Stiennon et al., 2020), but these  
073 require extensive training data and may reduce model capa-  
074 bilities.

### 075 2.2. Programmable Dialog Management

076 The concept of programmable dialog management has  
077 emerged as a promising direction for LLM security. Unlike  
078 static rule-based systems, programmable frameworks allow  
079 security policies to be expressed as executable code that can  
080 adapt to context. NeMo Guardrails represents a significant  
081 advancement in this direction, providing a domain-specific  
082 language (Colang) for expressing security logic.

## 083 3. NeMo Guardrails Architecture

### 084 3.1. Event-Driven Defense Layer

085 NeMo Guardrails operates as an event-driven conversation  
086 manager, fundamentally different from traditional Web Ap-  
087 plication Firewalls (WAFs). While WAFs operate at the  
088 network or application layer, Guardrails operates at the  
089 semantic level, auditing and intervening in every state trans-  
090 sition during interactions.

091 The core of the Guardrails runtime is an event loop that  
092 continuously processes events and generates new events.  
093 This design enables the system to handle complex conversa-  
094 tional scenarios in a non-linear manner. The processing  
095 flow can be decomposed into three critical control stages,  
096 each constituting part of a defense-in-depth strategy.

#### 103 3.1.1. STAGE 1: USER INTENT CANONICALIZATION

104 The first principle in security for handling unstructured data  
105 is normalization. User utterances contain variants, noise,  
106 and potentially adversarial perturbations. If security rules  
107 are written directly against raw text, the rule base will ex-

108 pand infinitely and be easily bypassed.

**Mechanism:** The runtime first triggers the generate\_user\_intent action. Rather than relying on hardcoded rules, this action uses vector search to find the most similar intent examples (few-shot examples) in the configuration, then prompts the LLM to classify the current input into a predefined “canonical form.” For example, whether the user inputs “hey,” “hello,” or “good morning,” the system normalizes it to the express greeting event.

**Security Value:** This layer serves as the first line of defense. By forcing inputs to map to a finite set of intents, prompt injection attacks attempting to use homophone substitution, Unicode obfuscation, or similar techniques often fail because they cannot match valid intents, or are forced into the unhandled category, triggering default security fallback policies.

#### 3.1.2. STAGE 2: NEXT-STEP PREDICTION AND FLOW DECISION

Once intent is determined, the system must decide the next action. This is where Colang flows come into play.

#### Hybrid Control Model:

- **Deterministic Path:** If a Colang script explicitly defines a flow check security that mandates bot refuse when ask about sensitive data is detected, the runtime will enforce this path, completely bypassing the LLM’s generation logic. This is the foundation for implementing “hard constraints.”
- **Probabilistic Path:** If no matching explicit flow exists, the system falls back to the LLM, requesting it to predict the next logical step. This design allows the system to maintain flexibility when handling undefined security boundaries while enforcing absolute control within security boundaries.

**Architectural Insight:** This hybrid model resolves the contradiction between overly rigid traditional rule engines and uncontrollable pure LLM systems. Security teams can write high-priority “security flows” to cover all high-risk scenarios (e.g., politically sensitive topics, competitor mentions), while leaving low-risk casual conversation to the model’s free generation.

#### 3.1.3. STAGE 3: BOT UTTERANCE GENERATION AND OUTPUT AUDITING

Even after the system determines what the bot “should” say (e.g., bot express apology), the specific generated text still requires LLM completion, introducing secondary risk.

Table 1. Guardrail Categories and Their Functions

Category	Stage	Core Function
Input Rails	Pre-processing	Intercept before input enters dialog manager. Lowest computational cost defense layer.
Dialog Rails	State management	Control dialog flow and state. Core logic based on definitions.
Retrieval Rails	RAG integration	Filter and sanitize retrieved chunks in RAG pipeline.
Execution Rails	Tool calling	Validate LLM parameters and return values for external tool calls (APIs, databases).
Output Rails	Post-processing	Validate final generated text.

**Mechanism:** The `generate_bot_message` action is responsible for generating the final text.

**Output Rails:** Before text is returned to the user, output rails are triggered. This is not merely simple keyword filtering, but can invoke specialized models (such as Llama Guard or custom PII detection models) for semantic scanning of generated content. If hallucinations or sensitive information leakage are detected, the system can intercept the message and replace it with a preset safe response.

### 3.2. Five-Category Guardrail System

NeMo Guardrails subdivides security logic into five categories, constructing a comprehensive defense matrix. This classification not only facilitates logical decoupling but also enables experts from different domains (e.g., compliance experts, security engineers) to maintain their respective rule sets.

## 4. Colang Language Mechanism

### 4.1. From Rule Engine to Asynchronous State Machine

Colang is the soul of NeMo Guardrails, defining how security policies are expressed. The evolution from Colang 1.0 to 2.0 represents a fundamental change in the underlying runtime architecture, directly affecting system concurrency performance and expressiveness.

### 4.2. Colang 1.0 Limitations: Synchrony and Rigidity

Early Colang 1.0 design was closer to traditional Intent-Slot filling systems. Its primary architectural bottleneck was **synchronous blocking execution**.

**State Explosion:** In 1.0, handling complex multi-turn conversations often led to combinatorial explosion in flow definitions. Any possible context switch required explicit definition, making codebase difficult to maintain.

**Performance Bottlenecks:** Actions in 1.0 were blocking. If an input rail needed to call a remote PII detection API,

the entire conversation thread would be suspended, causing significant latency accumulation. This is unacceptable when handling high-concurrency enterprise applications.

content doesn't contain unauthorized sensitive data.

Prevent "data poisoning" at runtime.

Colang 2.0 is a refactoring for the generative AI era, introducing Python-like syntax structures and asynchronous primitives, transforming Guardrails into a high-performance concurrent state machine.

### 4.3. Colang 2.0: Pythonic Asynchronous Concurrent Runtime

This is 2.0's greatest contribution to security architecture. The introduction of `await` and `start` keywords enables parallel security detection.

**Parallel Rails Pattern:** In security architecture, we pursue "zero trust," meaning we must perform multi-dimensional checks on every interaction (toxicity, PII, hallucinations, jailbreaking). In Colang 1.0, these checks were serial ( $T_{\text{total}} = T_{\text{toxic}} + T_{\text{pii}} + T_{\text{jailbreak}}$ ). In 2.0, we can concurrently launch these detection tasks using the `start` keyword, with the main flow waiting for all tasks to complete via `await`.

**Latency Optimization:** This makes total latency approximately equal to the slowest detection service time ( $T_{\text{total}} \approx \max(T_{\text{toxic}}, T_{\text{pii}}, T_{\text{jailbreak}})$ ), dramatically reducing the performance cost of security compliance.

### 4.3.2. EXPLICIT ENTRY POINTS AND GENERATION OPERATORS

Colang 2.0 introduces an explicit `main` flow as the entry point, eliminating the ambiguity of flow trigger conditions in 1.0 and enhancing system predictability.

**Generation Operator (...):** This is an extremely important syntactic sugar. It allows developers to explicitly delineate the boundary between "deterministic logic" and "generative logic" in Colang scripts. For example, `answer = ... "Generate a response based on X"`. For auditing purposes, this makes code review clear: all non-...

165 parts are hardcoded, trusted security logic, while . . . -  
 166 marked sections are LLM generation regions requiring fo-  
 167 cused monitoring.

#### 169 4.3.3. STANDARD LIBRARY AND MODULARITY

170 2.0 introduces a standard library (CSL) and import mecha-  
 171 nism. This means security teams can develop a standardized  
 172 `corporate-security-policy.co` library containing all mandatory  
 173 input-output check logic. Development  
 174 teams from various business lines must import this core li-  
 175 brary when building their bots. This architecturally achieves  
 176 **centralized control and distributed execution** of security  
 177 policies.

#### 179 4.4. State Management and Context Awareness

181 Colang 2.0 enhances variable scope and lifecycle manage-  
 182 ment. Security policies can now depend on complex context  
 183 variables.

185 **Accumulative Risk Scoring:** We can define a global vari-  
 186 able `$user_risk_score`. Each time a user triggers a  
 187 minor violation (e.g., using profanity), the flow logic can  
 188 increment this score. When the score exceeds a threshold,  
 189 a circuit breaker mechanism terminates the session. This  
 190 **state-accumulation**-based defense strategy is far more pow-  
 191 erful than single-shot stateless filters, effectively defending  
 192 against attackers' "crescendo attacks" conducted through  
 193 multi-turn conversations.

### 195 5. NVIDIA NIM Integration: Zero-Trust 196 Infrastructure

198 If NeMo Guardrails is responsible for application-layer log-  
 199 ical security, then NVIDIA NIM (NVIDIA Inference Mi-  
 200 croservices) is responsible for underlying computational  
 201 and supply chain security. As security architects, we  
 202 cannot assume models run in a security vacuum; infrastructure  
 203 integrity is the foundation of trust.

#### 205 5.1. Container Supply Chain Security

206 NIM consists of hardened and signed container images  
 207 containing model weights and inference engines (e.g.,  
 208 TensorRT-LLM).

210 **Artifact Signing and Verification:** Every NIM image is  
 211 cryptographically signed by NVIDIA before release. When  
 212 deploying in Kubernetes clusters, we can configure admis-  
 213 sion controllers to enforce image signature verification. This  
 214 eliminates "supply chain poisoning" risk, where attackers  
 215 replace images and implant backdoors in models.

216 **Safetensors Format:** NIM prioritizes using the safetensors  
 217 format rather than traditional Python pickle serialization for

loading model weights. Pickle has well-known deserialization  
 vulnerabilities allowing arbitrary code execution during  
 model loading. Using safetensors fundamentally eliminates  
 this class of serious Remote Code Execution (RCE) risks.

**VEX and Vulnerability Management:** NVIDIA provides  
 VEX (Vulnerability Exploitability eXchange) records. In  
 enterprise environments, scanners often report numerous  
 base image CVEs. VEX records can explicitly indicate  
 which CVEs are unexploitable under NIM's specific config-  
 uration, significantly reducing Security Operations Center  
 "alert fatigue" and allowing teams to focus on real threats.

### 5.2. Inter-Service Communication Security: mTLS and Service Mesh

In production environments, Guardrails services typically act as gateways, with backend connections to multiple NIM services (e.g., Main NIM for generation and Safety NIM for detection). Communication between these services must follow zero-trust principles.

**mTLS Encryption:** Through integration with service meshes like Istio or Linkerd, all communication between Guardrails and NIM uses mutual TLS (mTLS) encryption. This not only prevents Man-in-the-Middle (MITM) attacks but, more importantly, provides strong identity authentication.

**Identity Verification:** mTLS ensures that Guardrails services can only connect to legitimate NIM instances holding valid certificates, preventing cluster-internal attackers from launching malicious Pods masquerading as model services to steal users' sensitive prompt data (which may contain PII).

**Deployment Pattern:** We recommend the **Gateway/Service Pattern** rather than Sidecar pattern for deploying Guardrails. The gateway pattern allows centralized policy enforcement, where one Guardrails instance can route and protect multiple backend model services, facilitating independent scaling and security policy upgrades.

## 6. Advanced RAG Security and Fact-Checking

Retrieval-Augmented Generation (RAG) is the mainstream pattern for enterprise applications but introduces specific security risks such as hallucinations and misattribution. NeMo Guardrails provides an evidence-based verification mechanism.

### 6.1. NLI-Based Fact-Checking

NeMo's `check_facts` action is not simple keyword matching but advanced verification logic based on Natural Language Inference (NLI) or using LLM-as-a-Judge.

220  
221  
222  
223  
224  
225  
226  
227  
228  
229  
230**Workflow:**

1. The system retrieves relevant knowledge chunks, stored in the \$relevant\_chunks variable.
2. The LLM generates an answer.
3. The check\_facts flow is triggered. This flow constructs a verification prompt, using \$relevant\_chunks as the "premise" and the generated answer as the "hypothesis."
4. The model judges whether the "premise" entails the information in the "hypothesis."

**Citation Verification Logic:** Although NeMo itself doesn't directly provide "citation formatting" magic functionality, through the above entailment relationship check, it essentially achieves citation authenticity verification. If the model generates a claim but cannot find evidence in \$relevant\_chunks, the entailment check will fail, and the guardrail will intercept the answer. This effectively prevents the model from fabricating facts or citing non-existent sources.

**6.2. SelfCheckGPT Hallucination Detection Algorithm**

For non-RAG scenarios or to enhance RAG robustness, NeMo implements an algorithm similar to SelfCheckGPT.

**Random Sampling Consistency:** The core idea of this algorithm is that if a model's answer to the same prompt is factual, multiple sampling results should be consistent at the factual level; if it's a hallucination, multiple sampling results often diverge.

**Implementation:** The system generates  $N$  additional sample answers at high temperature in the background, then compares the consistency between the main answer and these samples. If consistency is low, the system judges it as a hallucination and intercepts or flags it. This is a self-consistency check that doesn't rely on external knowledge bases. Although computationally expensive, it is extremely valuable in high-risk scenarios.

**7. Data Security and NeMo Curator**

Security doesn't only exist at inference time but begins at the data level. If RAG knowledge bases are contaminated or fine-tuning data contains PII, inference-time guardrails will face enormous pressure. NeMo Curator is a critical tool for ensuring data supply chain security.

**7.1. PII Identification and Sanitization**

NeMo Curator employs a hybrid strategy for PII cleaning to balance efficiency and accuracy.

**Rule Engine (Presidio):** For format-fixed PII (e.g., ID numbers, emails, phone numbers), Presidio is used for high-speed scanning based on regex and logic.

**LLM Semantic Identification:** For unstructured, context-dependent PII (e.g., "the manager sitting by the window" or implicit home address descriptions), Curator uses large models like Llama 3 70B for semantic-level identification. According to NVIDIA testing, this LLM-based method achieves 26% higher recall on core PII categories compared to pure rule-based methods.

**Sanitization Strategy:** Supports `redact` (direct deletion) and `replace` (replace with placeholders like `\{\{PERSON\}\}`). We recommend using the replacement strategy, as it preserves sentence grammatical structure, benefiting subsequent model fine-tuning or RAG retrieval effectiveness.

**7.2. Large-Scale Distributed Processing**

Curator is built on Dask and RAPIDS, supporting distributed processing across multi-GPU nodes. This means it can rapidly clean enterprise data lakes at petabyte scale. This is crucial for preventing "data leakage" in RAG systems—ensuring retrieved chunks are clean and PII-free before entering prompts.

**8. Performance Engineering: Addressing the "Latency Tax"**

Introducing Guardrails inevitably brings latency. In architectural design, we must minimize this "latency tax" through technical means to meet real-time interaction requirements.

**8.1. Streaming Architecture and Chunked Verification**

Traditional synchronous verification requires all generated text to complete before detection, leading to extremely high Time-To-First-Token (TTFT). NeMo introduces a streaming verification mechanism.

**Stream First Strategy:** After configuring `stream_first: True`, LLM-generated tokens are immediately streamed to the client, minimizing user-perceived latency.

**Chunked Parallel Detection:** Simultaneously, the Guardrails service buffers streaming tokens into chunks (e.g., 128 tokens) in the background. Once the buffer is full, it immediately performs asynchronous security scanning on that chunk.

**Latency Cutoff:** If a chunk is detected as violating policies, Guardrails immediately cuts off the stream and sends a preset error message to override or append to the frontend. Although this may cause users to briefly see partial violat-

275 ing content (very short time window), it achieves the best  
 276 balance between user experience (extremely fast response)  
 277 and security.

278 **Chunk Size Trade-off:** Smaller chunks enable faster detection  
 279 but may lose context; larger chunks (e.g., 256 tokens)  
 280 provide more complete context (beneficial for hallucination  
 281 detection) but increase latency. We typically recommend  
 282 128 tokens as a balance point.

## 284 8.2. Latency Model Analysis

285 In the optimized architecture, the total latency model is:

$$286 \quad L_{\text{total}} = L_{\text{input\_rail}} + L_{\text{first\_token}} + \epsilon \quad (1)$$

287 where  $L_{\text{input\_rail}}$  is compressed through parallelization (parallel  
 288 execution), i.e., toxicity detection and PII detection  
 289 run in parallel, taking the maximum value rather than the sum.  
 290 Output rail latency is hidden in the generation process  
 291 through streaming mechanisms and no longer blocks  
 292 first-token display.

## 300 9. Security Evaluation and Red Teaming

301 As security architects, we cannot only build defenses but  
 302 must conduct adversarial testing. NVIDIA’s AI Red Team-  
 303 ing methodology provides evaluation standards.

### 306 9.1. Methodology: Limit-Seeking

307 NVIDIA’s red team testing emphasizes “limit-seeking, non-  
 308 malicious” exploration. This means the testing goal is not  
 309 simply to break but to probe model behavior boundaries.

310 **Tool Integration:** The NeMo ecosystem integrates auto-  
 311 mated scanning tools like Garak to probe known vulnerabil-  
 312 ity patterns (e.g., prefix injection, goal hijacking).

### 315 Specific Attack Vector Defense:

- 316 • **Jailbreaking:** Uses specialized JailbreakDetect NIM,  
 317 a model specifically trained against adversarial prompts  
 318 (e.g., DAN mode, Base64 encoded attacks) that can  
 319 identify structured attacks attempting to bypass secu-  
 320 rity policies.

- 321 • **Prompt Injection:** By strictly isolating user input in  
 322 specific regions of prompt templates and coordinating  
 323 with Input Rails for semantic intent recognition, NeMo  
 324 effectively reduces instruction injection risk.

## 10. Conclusion and Strategic Recommendations

In summary, the NVIDIA NeMo ecosystem constructs an enterprise-grade AI security loop through NeMo Guardrails’ programmability, Colang 2.0’s asynchronous concurrency capabilities, and NIM’s infrastructure security.

### Core Recommendations:

- Embrace Colang 2.0 Comprehensively:** Although 1.0 is still in use, immediate migration planning to 2.0 should begin. 2.0’s asynchronous features are the only path to solving multi-modal security detection latency bottlenecks.
- Implement Zero-Trust Networking:** Enforce mTLS between NIM and Guardrails. Do not expose model services on cluster networks.
- Emphasize Both Data and Models:** Deploy Guardrails not only at inference time but also use Curator for cleaning at the data level. Dirty data renders even the best guardrails ineffective.
- Hybrid Defense Strategy:** Leverage NeMo’s flexibility to combine deterministic flow logic (for high-risk red lines) and probabilistic LLM detection (for long-tail risks), constructing a multi-layer defense system.

The NeMo project demonstrates that AI security is no longer esoteric but a systems engineering discipline that can be solved through precise engineering architecture.

## Impact Statement

This paper presents work whose goal is to advance the field of Machine Learning security, specifically in the context of Large Language Model deployment. The security analysis and architectural insights presented here can help organizations deploy LLM-based systems more safely in enterprise environments. There are many potential societal consequences of our work, including improved safety of AI systems in critical applications, but also potential concerns about over-reliance on automated security mechanisms. We encourage further research into the limitations and failure modes of guardrail systems.

## References

Goodside, R. Prompt injection, 2021. Available at <https://simonwillison.net/2022/Sep/12/prompt-injection/>.

Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y., Dai, W., Yu, A. M., et al. Survey of halluci-

330 nation in natural language generation. *ACM Computing*  
331 *Surveys*, 55(12):1–38, 2023.

332  
333 Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C.,  
334 Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A.,  
335 et al. Training language models to follow instructions  
336 with human feedback. *Advances in Neural Information  
337 Processing Systems*, 35:27730–27744, 2022.

338 Stiennon, N., Ouyang, L., Wu, J., Ziegler, D. M., Lowe,  
339 R., Voss, C., Radford, A., Amodei, D., and Christiano,  
340 P. F. Learning to summarize with human feedback. *Ad-*  
341 *vances in Neural Information Processing Systems*, 33:  
342 3008–3021, 2020.

343 Wei, A., Haghtalab, N., and Steinhardt, J. Jailbroken:  
344 How does llm safety training fail? *arXiv preprint  
345 arXiv:2307.02483*, 2023.

346  
347 White, J., Fu, Q., Hays, S., Sandborn, M., Olea, C., Gilbert,  
348 H., Elnashar, A., Spencer-Smith, J., and Schmidt, D. C.  
349 A prompt pattern catalog to enhance prompt engineering  
350 with chatgpt. *arXiv preprint arXiv:2302.11382*, 2023.

351

352

353

354

355

356

357

358

359

360

361

362

363

364

365

366

367

368

369

370

371

372

373

374

375

376

377

378

379

380

381

382

383

384